

Intro to normalizing flows

(with Python examples)

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International (CC BY-NC-SA 4.0) license.

Kamila Zdybał

kamilazdybal.github.io, kamila.zdybal@gmail.com

Please feel free to contact me with any suggestions, corrections or comments.

Preface

Normalizing flows continuously transform one probability distribution into another. Their core concept is behind *generative models*. This is an illustrative tutorial providing a gentle introduction to normalizing flows with a bunch of practical Python examples.

Keywords

normalizing flow | probability density | Python

1 Transformations between random variables

We begin with a simple example of transforming one probability density into another which will provide you with the necessary intuition to move on towards understanding normalizing flows.

1.1 A simple example

Consider a random variable X sampled from a Gaussian normal distribution with a zero mean and a standard deviation equal to 1,

$$X \in \mathcal{N}(0, 1). \quad (1)$$

We would like to transform samples of X using some nonlinear, invertible and differentiable transformation, T . Hence, to go from $X \rightarrow Y$, we apply T such that

$$Y = T(X). \quad (2)$$

Equivalently, since T is invertible, we can also go from $Y \rightarrow X$ in the following way:

$$X = T^{-1}(Y). \quad (3)$$

An example transformation that we will use here is the exponential function $T(\bullet) = \exp(\bullet)$, whose inverse is $T^{-1}(\bullet) = \ln(\bullet)$. Fig. 1 shows histograms of X and Y from this example, where we initially generated 10,000 samples of X . You can see that the distribution of Y is no longer Gaussian normal – applying a nonlinear T changes the distribution of a random variable.

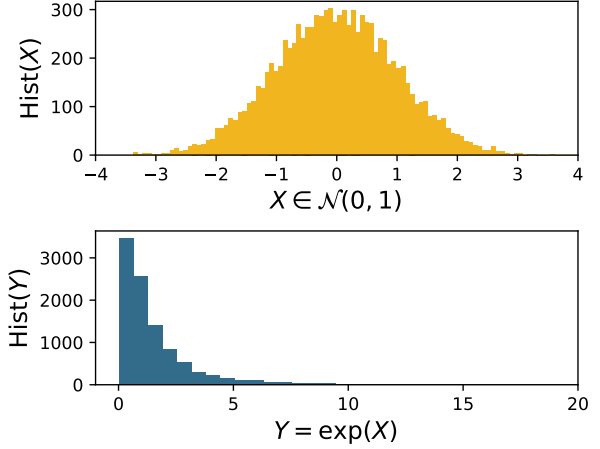


Figure 1: Histograms of two random variables, X and Y .

Let's denote the distribution of X as p_X and the distribution of Y as p_Y . We now pose the following question: Given p_X and T , what is p_Y ? It turns out that p_Y can be computed analytically using the change of variables formula:

$$p_Y(Y) = p_X(X) |\det J_T(X)|^{-1}, \quad (4)$$

where $J_T(X)$ is the Jacobian of the transformation T . Eq. (4) is equivalent to

$$p_Y(Y) = p_X(T^{-1}(Y)) |\det J_{T^{-1}}(Y)| \quad (5)$$

since $X = T^{-1}(Y)$ and, perhaps less trivially¹, $|\det J_{T^{-1}}(Y)| = |\det J_T(X)|^{-1}$. In the case of an exponential function, the Jacobian of the transformation can be computed as

$$J_T(X) = \frac{d(\exp(X))}{dX} = \exp(X). \quad (6)$$

Similarly, $J_{T^{-1}}(Y)$ is the Jacobian of the inverse transformation which in this case is

$$J_{T^{-1}}(Y) = \frac{d(\ln(Y))}{dY} = 1/Y. \quad (7)$$

¹We can show that $J_{T^{-1}}(Y) = J_T(X)^{-1}$ using the inverse function theorem. For an invertible and differentiable function T , we have by definition $T(T^{-1}(x)) = x$. By differentiating both sides with respect to x , we have $d/dx(T(T^{-1}(x))) = 1$, which, using the chain rule, becomes $dT/dx(T^{-1}(x)) \cdot dT^{-1}(x)/dx = 1$. Hence, the derivative of the inverse function is $dT^{-1}(x)/dx = \frac{1}{dT/dx(T^{-1}(x))}$, where we have also assumed that the derivative of T is non-zero for all x . Note, that if the derivative of T is zero, then T^{-1} is by definition not a function. Translating this to a Jacobian of the transformation T , we have: $J_{T^{-1}}(Y) = \frac{1}{J_T(T^{-1}(Y))} = \frac{1}{J_T(X)} = J_T(X)^{-1}$.




Figure 2: Probability density functions (PDFs) of two random variables, X and Y , where $X \in \mathcal{N}(0, 1)$ and $Y = \exp(X)$.

Fig. 2 shows probability densities for X and Y , where p_X is computed as

$$p_X(X) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(X-\mu)^2}{2\sigma^2}\right) \quad (8)$$

with $\sigma = 1$ and $\mu = 0$, and p_Y is computed either from Eq. (4) or Eq. (5).

1.2 Python computation

To compute p_Y using Eq. (4) in Python:

```
def p_transformed(X, p_X):
    p_Y = p_X(X, 0, 1) * np.abs(1/np.exp(X))
    return p_Y
```

and to compute p_Y using Eq. (5) in Python:

```
def p_transformed(Y, p_X):
    p_Y = p_X(T_inv(Y), 0, 1) * np.abs(1/Y)
    return p_Y
```

Both computations are equivalent, since $X = T^{-1}(Y)$ and $Y = \exp(X)$.